

# Image Warping and Alignment

# Basics

- A **digital** image – a version of the visual stimulus sampled at discrete locations, with discretized values
- Can be regarded as a function  $I = f(x,y)$  where  $(x,y)$  are spatial (integer) coordinates in typically a rectangular domain  $\Omega = [0,W-1] \times [0,H-1]$ .
- Each ordered pair  $(x,y)$  is called a **pixel**.
- The pixel is generally square (sometimes rectangular) in shape.

# Basics


- Pixel dimensions (height/width of the pixel) relate to the **spatial resolution** of the sensor in the camera that collects light reflected from a scene.
- Remember: the actual visual signal is analog, but digital cameras capture a **discrete** version of it, and also quantize the intensity values.
- That is they capture the visual stimulus only at specific points  $(x,y)$ , usually evenly spaced from each other in both X and Y directions.

# Basics

- In a typical photographic grayscale image, the intensity values  $f(x,y)$  lie in the range from 0 to 255 (8 bit integers).
- They are quantized versions continuous values corresponding to the actual light intensity that strikes a light sensor in the camera.

# Image Alignment

- Consider images  $I_1$  and  $I_2$  of a scene acquired through different viewpoints.



Pixels in **digital** correspondence (same coordinate values in the image domain  $\Omega$ , not necessarily containing measurements of the same physical point)

Pixels in **physical** correspondence (containing measurements of the same physical point, but not necessarily the same coordinate values in the image domain  $\Omega$ )

# Image Alignment

- $I_1$  and  $I_2$  are said to be aligned if for every  $(x,y)$  in the domain  $\Omega$ , the pixels at  $(x,y)$  in  $I_1$  and  $I_2$  are in physical correspondence.
- If not, the images are said to be **misaligned** with respect to each other.
- Or we say there is **relative motion** between the images.
- Image alignment (also called **registration**) is the process of correcting for the relative motion between  $I_1$  and  $I_2$ .





Image taken from the Brainweb database of the Montreal Neurological Institute

# Motion Models

- Let us denote the coordinates in  $I_1$  as  $(x_1, y_1)$  and those in  $I_2$  as  $(x_2, y_2)$ .

- Translation:  $\forall (x_1, y_1) \in \Omega, x_2 = x_1 + t_x, y_2 = y_1 + t_y$

$$\rightarrow \begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix}$$

- Rotation about point (o,o) anti-clockwise through angle  $\theta$

$$x_2 = x_1 \cos \theta - y_1 \sin \theta$$

$$y_2 = x_1 \sin \theta + y_1 \cos \theta$$

$$\begin{pmatrix} x_2 \\ y_2 \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} x_1 \\ y_1 \end{pmatrix}$$

2D Rotation matrix  
(orthonormal matrix)

# Motion Models

- Rotation about point  $(x_c, y_c)$  anti-clockwise through angle  $\theta$

$$x_2 = (x_1 - x_c) \cos \theta - (y_1 - y_c) \sin \theta + x_c$$

$$y_2 = (x_1 - x_c) \sin \theta + (y_1 - y_c) \cos \theta + y_c$$

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta & x_c \\ \sin \theta & \cos \theta & y_c \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 - x_c \\ y_1 - y_c \\ 1 \end{pmatrix}$$

- Perform translation such that  $(x_c, y_c)$  coincides with the origin  $(0,0)$ .
- Rotate about the new origin.
- Translate back.
- The extra ones (third row) are called **homogeneous coordinates** – they facilitate using matrix multiplication to represent translations.



# Motion Models

## ■ Rotation and translation:

$$x_2 = (x_1 - x_c) \cos \theta - (y_1 - y_c) \sin \theta + t_x$$

$$y_2 = (x_1 - x_c) \sin \theta + (y_1 - y_c) \cos \theta + t_y$$

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} = \begin{pmatrix} \cos \theta & -\sin \theta & t_x \\ \sin \theta & \cos \theta & t_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 - x_c \\ y_1 - y_c \\ 1 \end{pmatrix}$$

## ■ Affine transformation: (rotation, scaling and shearing) besides translation

Assumption: the  $2 \times 2$  sub-matrix  $\mathbf{A}$  is NOT rank deficient, otherwise it will transform two-dimensional figures into a line or a point

$$\begin{pmatrix} x_2 \\ y_2 \\ 1 \end{pmatrix} = \begin{pmatrix} A_{11} & A_{12} & t_x \\ A_{21} & A_{22} & t_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ y_1 \\ 1 \end{pmatrix}$$

# Motion Models


## ■ Scaling about the origin

Identity

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$x = v$$

$$y = w$$




Scaling

$$\begin{bmatrix} c_x & 0 & 0 \\ 0 & c_y & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$x = c_x v$$

$$y = c_y w$$



# Motion Models


## ■ Shearing:

Shear (vertical)

$$\begin{bmatrix} 1 & 0 & 0 \\ s_v & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$x = v + s_v w$$

$$y = w$$




Shear (horizontal)

$$\begin{bmatrix} 1 & s_h & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$x = v$$

$$y = s_h v + w$$



# Motion Models

- The 2D affine motion model (including translation in X and Y direction) includes 6 degrees of freedom.
- Note: this motion model accounts for in-plane motion only (example: not an appropriate model for “3D head profile view versus head frontal view”)
- Note: even with in-plane motion, there exist more complicated motion models, but we will stick to this one for now.


# Motion Models

- Composition of multiple types of motion is given by the multiplication of their corresponding matrices.
- So, if you first scale (matrix **S**) and then rotate (matrix **R**), then the resultant transformation = **RS**.
- Note: most motion compositions are not commutative (**RS** is not equal to **SR**).

# Motion Models

- In actual coding, you will typically not use matrices.
- Rather you will implement the formula as is.
- So why is matrix-based motion representation useful?
- Because it allows for a compact way to represent the composition of many different types of motion.

# Alignment with control points



Solve for unknown parameters using least-squares framework (i.e. pseudo-inverse)

Apply the motion based on these parameters to the first image

Control points: pairs of physically corresponding points – maybe marked out manually, or automatically using geometric properties of the image.

Number of control points  $k$  MUST be  $\geq u/2$ , where  $u$  = number of unknown parameters in the motion model (each point has two coordinates –  $x$  and  $y$ ). The number of control points is several times smaller than the number of image pixels.

$$\begin{pmatrix} x_{21} & x_{22} & \cdot & \cdot & x_{2k} \\ y_{21} & y_{22} & \cdot & \cdot & y_{2k} \\ 1 & 1 & \cdot & \cdot & 1 \end{pmatrix}$$

$$= \begin{pmatrix} A_{11} & A_{12} & t_x \\ A_{21} & A_{22} & t_y \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_{11} & x_{12} & \cdot & \cdot & x_{1k} \\ y_{11} & y_{12} & \cdot & \cdot & y_{1k} \\ 1 & 1 & \cdot & \cdot & 1 \end{pmatrix}$$

# Alignment with control points

- Not always feasible, if it requires manual intervention
- Error-prone
- There are methods for finding matching control points automatically (eg: the popular SIFT technique), but we will not cover them in this course.



# Alignment with mean squared error

- Mean squared error is given by:

$$MSSD = \frac{1}{N} \sum_{x,y \in \Omega} (I_1(x,y) - I_2(x,y))^2, N = \# \text{pixels in field of view (see defn. in later slides)}$$

- Find motion parameters as follows:

$$\mathbf{T}^* = \arg \min T MSSD_T(I_1(\mathbf{v}), I_2(\mathbf{T}\mathbf{v}))$$

Find transformation matrix  $T$  which produces the least value of MSSD

$$\mathbf{T} = \begin{pmatrix} A_{11} & A_{12} & t_x \\ A_{21} & A_{22} & t_y \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{v} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

# Alignment with mean squared error

- For simplicity, assume there was only rotation and translation.
- Then we have

$$\mathbf{T}^* = \arg \min \text{MSSD}_T(I_1(\mathbf{v}), I_2(T\mathbf{v}))$$

$$\mathbf{T} = \begin{pmatrix} \cos \theta & -\sin \theta & t_x \\ \sin \theta & \cos \theta & t_y \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{v} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$


# Alignment with mean-squared error

- There are many ways to do this minimization. The simplest but least efficient way is to do a brute-force search.
- Sample  $\theta$ ,  $t_x$  and  $t_y$  uniformly from a certain range (example:  $\theta$  from  $-45$  to  $+45$ ,  $t_x$  or  $t_y$  from  $-30$  to  $+30$ ).
- Apply this motion to  $I_1$  keeping  $I_2$  fixed (alternatively to  $I_2$  keeping  $I_1$  fixed, if the matrix is invertible), and compute the MSSD.
- Each time, compute the MSSD. Pick the parameter values corresponding to minimum MSSD.

# Image Warping


## ■ Forward warping method:

- ✓ Apply the transformation  $\mathbf{T}$  to every coordinate vector  $\mathbf{v} = [x \ y]$  in the original image to yield new coordinate  $\mathbf{T}\mathbf{v}$ .
- ✓ Copy the intensity value from  $\mathbf{v}$  in the original image to the new location in the warped image. Careful handling is needed if  $\mathbf{T}\mathbf{v}$  is not an integer (highly likely).



# Image Warping


- Forward warping:
  - Can leave the destination image with some holes if you scale up.
  - Can lead to multiple answers in one pixel if you scale down.




# Image warping

## ■ Reverse warping:

- ✓ For every coordinate  $\mathbf{v} = [x \ y]$  in the destination image, copy the intensity value from coordinate  $\mathbf{T}^{-1}\mathbf{v}$  in the original image.
- ✓ In case of non-integer value, perform interpolation (nearest neighbor or bilinear)



# Interpolation



## Nearest neighbor method:


- Use value  $a_4$  (as the pixel that is nearest to the red point contains value  $a_4$ )

## Bilinear method:


- Use the following value, a weighted combination of the four neighboring pixel values, with more weight to nearer values:

$$(Ba_4 + Aa_3 + Ca_1 + Da_2) / (A + B + C + D) = (Ba_4 + Aa_3 + Ca_1 + Da_2)$$

as  $A + B + C + D = 1$  for unit area pixels



# Bilinear interpolation in more detail



We interpolate first in the X direction:

$$f(x, y_1) = \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21})$$

$$f(x, y_2) = \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22})$$

We then interpolate first in the Y direction:



$$f(x, y) = \frac{y_2 - y}{y_2 - y_1} f(x, y_1) + \frac{y - y_1}{y_2 - y_1} f(x, y_2)$$

$$= \frac{y_2 - y}{y_2 - y_1} \left( \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21}) \right) + \frac{y - y_1}{y_2 - y_1} \left( \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22}) \right)$$

In both cases, notice that higher weight is given to the pixels that are closer to P.



# Bilinear interpolation in more detail



$$f(x, y) = \frac{y_2 - y}{y_2 - y_1} \left( \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21}) \right) + \frac{y - y_1}{y_2 - y_1} \left( \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22}) \right)$$

$$= \frac{y_2 - y}{y_2 - y_1} \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{y_2 - y}{y_2 - y_1} \frac{x - x_1}{x_2 - x_1} f(Q_{21}) + \frac{y - y_1}{y_2 - y_1} \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{y - y_1}{y_2 - y_1} \frac{x - x_1}{x_2 - x_1} f(Q_{22})$$

This formula will remain unchanged if you first interpolated in the Y direction and then in the X direction. Verify this yourself.

# Bilinear interpolation in more detail

- Here we are approximating the image intensity in the form of the following bilinear function:

$$f(x, y) = a_0 + a_1x + a_2y + a_3xy$$

- Here  $a_0, a_1, a_2, a_3$  are scalar coefficients. The value of  $f(x, y)$  is known at the four corners of a pixel.
- The function would have been linear if the term in  $xy$  were absent (i.e. if  $a_3 = 0$ ).

# Bilinear interpolation in more detail

- How do we determine the coefficients  $a_0, a_1, a_2, a_3$ ?

$$\begin{pmatrix} 1 & x_1 & y_1 & x_1 y_1 \\ 1 & x_1 & y_2 & x_1 y_2 \\ 1 & x_2 & y_1 & x_2 y_1 \\ 1 & x_2 & y_2 & x_2 y_2 \end{pmatrix} \begin{pmatrix} a_0 \\ a_1 \\ a_2 \\ a_3 \end{pmatrix} = \begin{pmatrix} f(x_1, y_1) \\ f(x_1, y_2) \\ f(x_2, y_1) \\ f(x_2, y_2) \end{pmatrix} = \begin{pmatrix} f(Q_{11}) \\ f(Q_{12}) \\ f(Q_{21}) \\ f(Q_{22}) \end{pmatrix}$$


- These coefficients can be obtained by inverting the  $4 \times 4$  matrix.
- It can be shown (through tedious calculations) that the result of this is equivalent to the formula we derived earlier.

# Alignment with mean squared error


- In the ideal case, the MSSD between two perfectly aligned images is 0. In practice, it will have some small non-zero value even under more or less perfect alignment due to sensor noise or slight mismatch in pixel grids.

# Careful: field of view issues!










Fixed image (also called reference image)




Region of overlap (also called **field of view**) when the moving image is warped



Note: compute MSSD only over region of overlap.



Change in region of overlap (also called **field of view**), as the moving image is warped. The field of view in any of these figures consists of all those pixels not marked black. It represents the set of pixels  $(x,y)$  where both  $I_1(x,y)$  and  $I_2(x,y)$  are well defined.




# Alignment with mean squared error


- MSSD is called an “image similarity metric”.
- MAJOR ASSUMPTION: Physically corresponding pixels have the same intensity!

# Image alignment: Intensity changes in images


- Images acquired by different sensors (MR and CT, camera with and without flash, etc.)
- Changes in lighting condition



MR-PD



MR-T<sub>1</sub>






# Image alignment: Intensity changes in images

- If following relationship (function  $g$ ) exists and we **knew** it, the solution is easy:

$$I_1(x_1, y_1) = g(I_2(x_2, y_2)), \forall (x_1, y_1), (x_2, y_2) \in \Omega$$



Physically corresponding points

$$\text{transformed - MSSD} = \frac{1}{N} \sum_{x,y \in \Omega} (g(I_2(x, y)) - I_1(x, y))^2$$

$N$  = # pixels in the field of view (see earlier for definition of field of view)

# Image alignment: Intensity changes in images

- What if the relationship exists in the following linear form, but we knew it only partially?

$$I_1(x_1, y_1) = aI_2(x_2, y_2) + b, \forall (x_1, y_1), (x_2, y_2) \in \Omega$$

Physically corresponding points

$$NCC = \left| \frac{\sum_{(x,y) \in \Omega} (I_1(x,y) - \bar{I}_1)(I_2(x,y) - \bar{I}_2)}{\sqrt{\sum_{(x,y) \in \Omega} (I_1(x,y) - \bar{I}_1)^2 \sum_{(x,y) \in \Omega} (I_2(x,y) - \bar{I}_2)^2}} \right|$$

Normalized cross-correlation, also called correlation-coefficient

$\bar{I}_1, \bar{I}_2$  : average value of images  $I_1, I_2$

We are taking the absolute value here, to take care of cases where one image has positive values and the other has negative values. We are assuming a linear relationship between the intensities of  $I_1$  and  $I_2$  but we do not assume knowledge of the scalar coefficients  $a$  and  $b$ .

# Image alignment: Intensity changes in images

$$NCC = \left| \frac{\sum_{(x,y) \in \Omega} (I_1(x,y) - \bar{I}_1)(I_2(x,y) - \bar{I}_2)}{\sqrt{\sum_{(x,y) \in \Omega} (I_1(x,y) - \bar{I}_1)^2 \sum_{(x,y) \in \Omega} (I_2(x,y) - \bar{I}_2)^2}} \right|$$

$\bar{I}_1, \bar{I}_2$  : average value of images  $I_1, I_2$

Normalized cross-correlation, also called correlation-coefficient. The NCC is like the absolute normalized dot product between mean-deducted images.

$$\mathbf{T}^* = \arg \max NCC_T(I_1(\mathbf{v}), I_2(T\mathbf{v}))$$



$$\mathbf{T} = \begin{pmatrix} \cos \theta & -\sin \theta & t_x \\ \sin \theta & \cos \theta & t_y \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{v} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

# Image alignment: intensity changes in images?

- Assume there exists a functional relationship between intensities at physically corresponding locations in the two images.
- But suppose we didn't know it (most practical scenario) and couldn't find it out.
- We will use image histograms!

# Image Histogram


- In a typical digital image, the intensity levels lie in the range  $[0, L-1]$ .
- The (normalized) histogram of the image is a discrete function of the form  $p(r_k) = n_k / HW$ , where  $r_k$  is the  $k$ -th intensity value, and  $n_k$  is the number of pixels with that intensity. ( $H, W$  = image dimensions)
- Sometimes, we may consider a range of intensity values for one entry in the histogram, in which case  $r_k = [r_{k'}^{\min}, r_{k'}^{\max}]$  represents an intensity bin, and  $n_k$  is the number of pixels whose intensity falls within this bin.
- Note  $p(r_k) \geq 0$  always, and all the  $p(r_k)$  values sum up to 1.




These are unnormalized histograms of two different images. That is, the entries of the histogram represent frequencies in this case without division by  $HW = \text{\#pixels}$ .

# Joint Image Histogram

- A joint image histogram is a function of the form  $p(r_{k_1}, r_{k_2})$  where  $r_{k_1}$  and  $r_{k_2}$  represent intensity bins from the two images  $I_1$  and  $I_2$  respectively.
- $p(r_{k_1}, r_{k_2})$  = number of pixels  $(x, y)$  such that  $I_1(x, y)$  and  $I_2(x, y)$  lie in bins  $r_{k_1}$  and  $r_{k_2}$  respectively, divided by HW.




$I_2$




Registered images: joint histogram plot looks "streamlined"





Values  
in  $I_2$   
from 0  
to 60



Values in  $I_1$  (from 0 to 60)

Misaligned images: joint histogram plot  
looks "dispersed"

We need a method to quantify how dispersed a  
joint histogram actually is.

# Measure of dispersion

- Consider a discrete random variable  $X$  with normalized histogram  $p(X)$  [also called the probability mass function].
- The entropy of  $X$  is a measure of the uncertainty in  $X$ , given by the following formula:

$$H(X) = - \sum_{x \in DX} p(X = x) \log_2 p(X = x)$$

$DX$  = discrete set of values that  $X$  can take

- Note: entropy is a function of the **normalized histogram** of  $X$ .
- **Not** a function of the **actual values** of  $X$ .
- The entropy is always non-negative.

# Entropy

- The entropy is maximum if  $X$  has a discrete uniform distribution, i.e.  $p(X=x_1) = p(X=x_2)$  for all values  $x_1$  and  $x_2$  in  $DX$ . The maximum value is  $\log(|DX|)$ .
- The entropy is minimum (zero) if the normalized histogram of  $X$  is a Kronecker delta function, i.e.  $p(X=x_1) = 1$  for some  $x_1$ , and  $p(X=x_2) = 0$  for all  $x_2 \neq x_1$ .

# Joint entropy

- The joint entropy of two random variables  $X$  and  $Y$  is given as follows:

$$H(X, Y) = - \sum_{x \in DX} \sum_{y \in DY} p(X = x, Y = y) \log_2 p(X = x, Y = y)$$

- Maximum entropy:

-Uniform distribution on  $X$  and  $Y$ : entropy value  $\log(|DX||DY|)$  where  $DX, DY$  are the set of discrete values that  $X$  and  $Y$  can acquire

- Minimum entropy:

-Kronecker delta (i.e. a PMF with all probability concentrated on only one entry with others being 0): entropy value 0 = non uncertainty

# Joint entropy

- Minimizing joint entropy is one method of aligning two images with different intensity profiles.


$$\mathbf{T}^* = \arg \min_{\mathbf{T}} H(I_1(\mathbf{v}), I_2(\mathbf{T}\mathbf{v}))$$

$$\mathbf{T} = \begin{pmatrix} A_{11} & A_{12} & t_x \\ A_{21} & A_{22} & t_y \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{v} = \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

$I_1$



$I_2$ : obtained by squaring the intensities of  $I_1$ , and rotating  $I_1$  anticlockwise by 30 degrees.



$I_2$  treated as moving image,  $I_1$  treated as fixed image. Joint entropy minimum occurs at -30 degrees.

# Components of an Image Alignment Algorithm

- Choice of a metric to optimize (joint entropy or mean squared error)
- Choice of a motion model (only translation, translation + rotation, affine, etc.)
- Choice of an interpolation algorithm to generate the warped image
- Choice of an optimization algorithm (here, we just used brute force search)


# Image Alignment: applications, related problems

- Template matching
- Image Panoramas



# Template Matching

- Look for the occurrence of a template (a smaller image) inside a larger image.
- Example: eyes within face image



Templates

# Template matching

- Let  $T$  = template image (smaller) of size  $h \times w$  and  $J$  = larger image of size  $H \times W$
- For every pixel  $(x,y)$  in  $J$ , consider a portion  $z_{xy} = J(x:x+w-1, y:y+h-1)$
- Select the portion  $z_{xy}$  with smallest MSSD compared to  $T$
- In some variants, the larger image  $J$  may be rotated with respect to  $T$ .
- In such cases, repeat the above procedure for every rotation of  $J$  from (say)  $-90$  to  $+90$  degrees.
- That is for every  $\theta$  from  $-90$  to  $+90$  degrees, let  $J_\theta$  be a rotated version of  $J$ . Now consider all  $z_{xy}$  in  $J_\theta$  and report the  $(x,y, \theta)$  triple that produces the least MSSD with respect to  $T$ .

# Image Panoramas



<http://cs.bath.ac.uk/brown/autostitch/autostitch.html>

A camera has a limited field of view. A scene may have much larger “area” than what can be captured from a single camera.

So one can acquire multiple images, each from a different viewpoint, and you need to stitch these together to form a panorama or mosaic. The stitching involves more complicated motion models (called homography) than what you have studied in this course.

# What we learnt..

- Affine motion model
- Forward and reverse image warping
- Field of view during image alignment
- Measures for Image alignment: sum of squared differences, normalized cross-correlation, joint entropy

# What we didn't learn

- Complicated motion models: higher degree polynomials, non-rigid models (example: motion of an amoeba, motion of the heart during the cardiac cycle, facial expressions, etc.)
- Efficient techniques for optimizing the measure for image alignment